

Tail-Robust Query Planning from Triple Storage

Huahai Yang
Juji, Inc.
San Jose, CA, USA
huahai.yang@gmail.com

ABSTRACT

Triple stores complicate cost-based query optimization because predicates on a single entity are correlated and relationship fanout depends on traversal direction. Datalevin [4] addresses these challenges with a layered strategy: index-maintained base-cardinality counts from nested entity-attribute-value (EAV) and attribute-value-entity (AVE) indexes, query-local reservoir sampling for conditional selectivity, and regularization toward durable storage statistics.

On a 277.9-million-triple instance of the 113-query Join Order Benchmark (JOB) [8], a reference run supplied the same optimizer with exact cardinalities (Exact). Full’s median paired execution-time ratio to Exact was $1.086\times$ (95% CI [1.044, 1.160]). In the layer ablations, each sampled condition used 30 paired seeds per query (3,390 trials). Shrinkage lowered Raw’s P99 from $12.26\times$ to $2.64\times$ (paired reduction $9.62\times$; 95% CI [2.69, 11.29]), while the observed count of Raw executions at or above $10\times$ fell from 48 to eight. Hiding predicate counts caused 30 timeouts and removing query-local sampling caused 80; Full had no timeouts in either experiment.

CCS CONCEPTS

• **Information systems** → **Query optimization**; *Database query processing*; *Database management system engines*.

KEYWORDS

Datalog, query optimization, cardinality estimation, sampling, robust query processing, entity-attribute-value storage

1 INTRODUCTION

Datalevin combines the declarative Datalog query language [10] with a flexible triple model. It stores facts as entity-attribute-value triples (E, A, V) rather than fixed-width rows, so applications can add attributes without schema migrations and traverse relationships in either direction.

EAV poses distinct optimization challenges. A *reference attribute* stores another entity’s ID as its value and therefore represents a relationship. A single logical entity spans many EAV triples, so queries over several entity types initially resemble self-joins, with reference attributes and ordinary value predicates interleaved. The relationship’s fanout—the number of output tuples per input—depends on whether traversal follows the forward or reverse order.

The optimizer must estimate three quantities: attribute-predicate selectivity, within-entity correlations, and directional relationship

fanout. In a row store, each histogram describes one typed column; in EAV, all domains share the V position and differ only by A , so a single V histogram would mix unrelated distributions. Per-attribute histograms avoid that mixing yet remain marginal: they cannot capture predicate co-occurrence or fanout conditional on the entities matched by earlier query steps. Precomputing statistics for open-ended attribute combinations and traversal directions is expensive and can lag the durable database. Triple-store systems accordingly replace one-dimensional histograms with RDF-specific synopses or live-index sampling [6, 12].

Datalevin instead derives cardinalities from its storage structure and query execution paths (Figure 2). Nested ordered indexes supply index-maintained base counts at $O(\log n)$ cost. Query-specific samples are drawn during planning from the entity ranges and traversal directions the query induces, not from a global precomputed sample; the optimizer obtains them by executing bounded versions of candidate scans and traversals.

An unlucky sample can cross a cost boundary where two join orders reverse. Datalevin combines the query-local sample with a durable fanout prior: shrinkage blends their means, while the production envelope takes their maximum. The central claim is compositional: counts, sampling, and uncertainty management address different dominant sources of execution-time tail, and removing any layer reintroduces severe failures.

This system experience paper makes three contributions:

- (1) A layered Datalog-planning architecture in which index-maintained counts resolve base cardinalities and query-local directional samples estimate correlation and fanout.
- (2) Two lightweight, storage-aware policies—prior shrinkage and a conservative envelope—that reduce the risk of sample-induced bad plans without offline learning or conventional histograms.
- (3) An ablation of counts, sampling, and regularization across all 113 JOB queries, with paired deterministic samples, plan-stability measurements, and execution-time tails, together with an exact-cardinality calibration of the fixed optimizer.

2 OPTIMIZER DESIGN

Datalevin stores triples in two orders. Entity-attribute-value (EAV) sorts by entity, then attribute, then value; this clusters each entity’s attributes together for efficient merge scans. Attribute-value-entity (AVE) sorts by attribute, then value, then entity, forming a secondary index for reverse-reference and value-equality lookups.

Illustrative query. Figure 1(a) shows a simplified JOB query in Datomic-style Datalog [5] that returns US-produced movie titles from 2005 to 2010. Each where clause pattern matches an (E, A, V) triple. Patterns sharing an entity variable form an *entity star*. Figure 1(b) shows these stars; collapsing them yields the optimizer graph in Figure 1(c), whose edges are reference patterns.

This paper is published under the Creative Commons Attribution 4.0 International (CC-BY 4.0) license. Authors reserve their rights to disseminate the work on their personal and corporate Web sites with the appropriate attribution, provided that you attribute the original work to the authors and CIDR 2027, 17th Annual Conference on Innovative Data Systems Research (CIDR ’27), January 24-27, Amsterdam, The Netherlands.




Figure 1: Illustrative query. (a) Datalog query. (b) Patterns grouped by entity variable into stars with attribute, value, and constraint spokes. (c) Collapsed query graph; stored references may be traversed in either direction.

On this entity-centered query graph, the optimizer applies cost-based dynamic programming. It estimates cardinalities for each candidate expansion and compares both indexed traversal directions, assembling a left-deep physical plan that joins one new star at a time to the accumulated prefix.

2.1 Entity stars and merge scans

Patterns sharing an entity variable are not independent relations—they are co-located in storage. In the example, the title and year patterns, together with the range constraint, form the ?t star. Datalevin’s merge scan walks candidate entity IDs in order and scans their requested attributes, applying predicates on the fly. Like the RDF pivot index scan [2], it produces one relation for the logical object rather than materializing one relation per attribute.

Star collapse turns dozens of triple patterns into a smaller graph whose nodes are stars and whose edges are references or shared values. The optimizer chooses a starting star, then an order and direction for crossing the edges; the example has three possible starts and two bidirectional edges.

2.2 Plan enumeration

The entity-star graph feeds plan enumeration. Datalevin applies Selinger-style dynamic programming over connected subsets of each query-graph component [11, 13]. It produces left-deep trees: each indexed expansion demands a base relation, so bushy plans are not considered. Disconnected components are planned independently, in parallel.

The enumerator preserves ordered keys and covers every connected ordered pair of stars, hence every valid start and first traversal. Once the candidate count for an n -star component exceeds $n(n-1)$, it retains only the cheapest plan per connected, unordered star set.

2.3 Physical alternatives

Enumeration fixes the traversal order, but the optimizer must also choose how each edge executes. The alternatives most relevant here are reference traversals. A forward reference feeds target entity IDs directly into a merge scan, whereas a reverse reference probes AVE to recover source entity IDs. These two physical directions implement the same logical edge but can differ sharply in cost and fanout; the estimator therefore treats them separately.

2.4 Cost model

Each candidate plan carries an estimated output cardinality and accumulated cost. Base cost scales with the number of candidates entering the initial index access and entity-star merge scan; pushed predicates and functions add fixed per-tuple factors. Extending a prefix through a reference adds $c_{\text{probe}}N_{\text{in}} + c_{\text{retrieve}}N_{\text{out}}$, while a hash join is proportional to the combined input sizes. These dimensionless operator weights rank plans; they do not predict wall-clock time. The optimizer chooses the lowest accumulated cost, using estimated output size to break ties.

Cardinality error cascades: N_{out} determines the current retrieval cost and becomes N_{in} for later steps. Underestimating a link can therefore make an expansion and its entire suffix appear cheap, displacing a selective filter to a later position. The three cardinality layers described in Section 3 target precisely these cost-model inputs.

3 MANAGING CARDINALITY UNCERTAINTY

3.1 Nested ordered indexes

EAV and AVE also provide cardinalities. A literal triple representation would repeat long shared prefixes, so Datalevin uses the LMDB DUPSORT facility [3], which maps one key to an ordered value list. This forms a nested B⁺ tree with page-level prefix compression. Internal nodes carry subtree entry counts, supporting $O(\log n)$ count and rank operations.

For an AVE point predicate $[?e : a \ v]$, the number of entities is the list count at key $(: a, v)$; rank metadata gives interval cardinality in logarithmic time. EAV similarly counts a known entity and attribute. For an attribute a , the durable prior $\mu_0(a) = |\{(e, a, v)\}| / |\{v : \exists e (e, a, v)\}|$ is the mean AVE postings-list length: expected source entities per distinct value. It is the prior used in Section 3.4 for AVE reverse-reference and value-equality paths; forward references use EAV and are costed separately. Unlike periodic histograms, these counts are updated with index entries in the same LMDB transaction; copy-on-write MVCC publishes counts and entries atomically to each reader snapshot, avoiding refresh lag. Datalevin also verifies a reported zero against the index before using it to prove emptiness.




Figure 2: The architecture for managing cardinality uncertainty. Counts resolve marginal cardinalities; samples estimate query-local correlation and directional fanout; the uncertainty policy regularizes those estimates before plan selection.

3.2 Predicate pushdown and base cardinalities

The compiler rewrites single-variable inequalities as index bounds, pushes eligible functions into attribute scans, and substitutes constants before planning. The count and scan therefore cover the same restricted range, avoiding a mismatch between estimated and executed predicates.

For a fixed attribute A , a bound clause fixes V to a constant; a free clause leaves V variable. Within each star, the optimizer uses the pattern with the smallest index-counted base cardinality for the initial scan and retains the alternatives. Small scans may execute during planning; large ones use a bounded sample.

3.3 Query-local directional sampling

Counts establish marginal selectivity, but do not reveal whether several attributes co-occur on the same entities or how many targets a relationship reaches. For a range of M candidates, Datalevin draws at most $m = 1,000$ ranks uniformly without replacement and retrieves only those entries [14]. Rank access costs $O(m \log M)$, or $O(\log M)$ in database size for fixed m .

Sampled entity IDs pass through the same merge scan and predicates that execution will use. The surviving fraction estimates the selectivity of the complete star, including local correlations.

Relationship estimates are directional. If a reference links A to B , forward expansion, reverse lookup, and value equality can have different fanouts. The estimator samples the candidate direction and access path. For n inputs, let x_i be the joined outputs from input i , including zero for no match. Directional join sampling produces the raw *link ratio*

$$\hat{r}_{A \rightarrow B}^{\text{raw}} = \frac{\text{sampled join outputs}}{\text{sampled inputs}} = \frac{1}{n} \sum_{i=1}^n x_i = \bar{x}.$$

This ratio predicts the multiplicative cardinality change along that physical direction: a value below 1 is selective, 1 is cardinality-preserving, and a value above 1 expands the input. The reverse ratio $\hat{r}_{B \rightarrow A}$ may be entirely different.

Sampling occurs at the base-plan stage; the optimizer caches star selectivities and adjacent link ratios for use in deeper prefixes, where a link maps N estimated rows to $N\hat{r}$. This assumes earlier joins do not select a different-fanout subpopulation; deeper resampling showed little development benefit.

Rare heavy hitters may be missed or overrepresented, and a noisy mean can put an estimate on the damaging side of a plan-cost crossover.

3.4 Storage-aware policies for sampling uncertainty

Given a sampled directional link ratio $\bar{x}$ and the attribute’s durable average fanout μ_0 , the unregularized baseline is

$$\hat{\mu}_{\text{raw}} = \max(\bar{x}, \mu_{\min}).$$

Here $\mu_{\min} = 1.0$ is a conservative optimizer floor: it never credits a sampled relationship traversal with reducing cardinality. This encodes the asymmetric cost of estimation errors—overestimation is usually tolerable, whereas underestimation can select a catastrophically bad join order. Raw preserves query-local evidence without accounting for sampling uncertainty. Against this baseline, we evaluate two policies that use the same reservoir sample but manage its uncertainty differently.

Shrinkage. The first policy moves the raw mean toward the durable average fanout:

$$\tilde{\mu} = \frac{n\bar{x} + k_{\text{eff}}\mu_0}{n + k_{\text{eff}}}, \quad k_{\text{eff}} = k(1 + \alpha \text{CV}^2). \quad (1)$$

Here k is the base strength of the storage prior, measured as an equivalent number of observations. CV^2 is the sample’s squared coefficient of variation, $\widehat{\text{Var}}(x)/\bar{x}^2$ for $\bar{x} > 0$ and zero otherwise; α controls how strongly dispersion increases the prior weight. Equation 1 gives the sample weight $n/(n + k_{\text{eff}})$ and the prior weight $k_{\text{eff}}/(n + k_{\text{eff}})$. A large, stable sample approaches $\bar{x}$, whereas a small or dispersed sample moves toward μ_0 . The final estimate is $\hat{\mu}_{\text{shrink}} = \max(\tilde{\mu}, \mu_{\min})$.

Conservative envelope. The second policy takes a one-sided envelope over the query-local observation and durable average:

$$\hat{\mu}_{\text{env}} = \max(\bar{x}, \mu_0, \mu_{\min}). \quad (2)$$

The envelope cannot fall below either source of evidence.

For input cardinality N , the relationship estimate $N\hat{\mu}$ propagates through later prefixes. The envelope is no smaller than shrinkage because it takes a maximum rather than an average. It is a decision rule, not an accuracy claim: underestimating an expansion can move it ahead of selective work, while overestimating it usually delays it. This may also postpone a beneficial traversal.

4 EVALUATION

We ablate counts, sampling, and fanout policy (Table 1). All conditions share the same plan space, physical alternatives, and operator costs; only the named cardinality evidence or fanout rule changes.

The primary outcome is catastrophic slowdown; P95 and P99 summarize slowdown relative to the corresponding Full-policy median. Full is the envelope condition. No counts withholds predicate-specific point and range counts from costing but retains μ_0 metadata and physical sizes needed for correct execution or sampling; equality falls back to μ_0 , other constraints to 10% of the attribute population. No sampling disables reservoir sampling of large base ranges; small ranges may still execute exactly, and unsampled links use μ_0 . Raw, Shrinkage, and Full receive identical sampled ranks within each query-seed group, so their comparison isolates how the fanout rule interprets the same evidence. A separate Exact reference, described below, calibrates Full without entering the layer ablation.

Table 1: Experimental conditions.

Condition	Predicate counts	Sampling	Fanout policy
Full	Used	Reservoir	Eq. 2
Shrinkage	Used	Reservoir	Eq. 1
Raw	Used	Reservoir	$\max(\bar{x}, \mu_{\min})$
No sampling	Used	None	Durable prior μ_0
No counts	Hidden	Reservoir	Eq. 2

4.1 Setup and protocol

JOB's 113 queries were designed to expose cardinality and join-order failures on a large, highly correlated Internet Movie Database (IMDb) snapshot [8]. We translated and verified the JOB workload in Datalevin Datalog; code and experiment artifacts are available online. The database contains 277,878,411 triples and occupies 17 GB.

Experiments ran on an Apple M3 Pro MacBook Pro with 36 GB memory and a 1 TB SSD, using OpenJDK 21.0.11 and Datalevin 1.0.0. On five development seeds, we compared candidate formulas and parameter settings and selected the best-performing tested configuration ($m = 1,000$, $k = 100$, $\alpha = 0.4$). We then fixed it, along with the timeout and analysis protocol, before the reported runs on disjoint seeds. A post-confirmatory one-factor sensitivity run used ten seeds on 59 queries whose Raw plans vary across seeds and on 20 stable controls, varying $m \in \{250, 1000, 4000\}$, $k \in \{0, 25, 100, 400\}$, and $\alpha \in \{0, 0.4, 1.6\}$. The frozen point ($m = 250$, $k = 400$, $\alpha = 1.6$) had no catastrophes; $m = 4000$ or weaker regularization produced four to nine (P99 8.56–10.04× versus the frozen 3.96×). This exploratory run did not retune the reported conditions.

Full, No counts, Raw, and Shrinkage use 30 deterministic passes. Within each pass, the harness shuffles all query-condition pairs with the pass seed, so conditions are interleaved rather than run in batches. It also derives one sample seed per query, shared across conditions. Because No sampling has one deterministic plan, nominal seeds would be pseudoreplicates; it instead uses ten randomized query-order schedules with paired Full controls. Sampled policies use each query's 30-run Full median; No sampling uses its 10-run control median.

Each query-condition pair gets a fresh plan cache and a 30-second timeout, charged at the limit. Queries run in a persistent, killable worker. After a worker start or timeout-induced restart, the

harness runs fixed query 6d and two unmeasured replays of the next scheduled query before recording its time. For every completed trial, it compares a result fingerprint—result cardinality plus a hash of the returned value—across conditions and repeated runs; all matched. Deterministic plan-only replays then use the same seeds to record operator and traversal order without contributing timing observations.

Against paired Full controls, No sampling cut median preparation from 124.6 to 5.3 ms, but raised the median preparation-plus-execution time among completed trials from 518 to 747 ms and timed out 80 times. We report these timing numbers here rather than with the plan enumeration because the ablation changes both the planning work and the plan it produces. A disjoint 100-seed plan-only census covers all 113 queries, preserving identical ranks within each query-seed group.

For query q , trial s , and policy p , slowdown is

$$S_{q,s,p} = \frac{T_{q,s,p}}{\text{median}_{s'} T_{q,s',\text{Full}}}. \quad (3)$$

Catastrophe definition. A trial is a catastrophe if it either reaches the 30-second timeout or its runtime is at least $10\times$ the median Full-policy runtime for the same query in the corresponding experiment ($S_{q,s,p} \geq 10$). Because the denominator is a Full median rather than each Full observation, an individual Full run can also cross this threshold. Reported 95% intervals use 2,000 paired-bootstrap replicates that resample queries and, within each resampled query, the seed-matched trials.

The 30-seed experiment has four queries (4c, 8c, 16b, and 17e) whose Full median exceeds 3 seconds; the separate 10-schedule controls for No sampling have seven (those four plus 10c, 13d, and 19d). For those query-experiment pairs, the $10\times$ threshold lies beyond the timeout, so catastrophe is equivalent to timeout. Table 2 therefore separates completed threshold crossings from timeouts. P95 and P99 remain query-relative: No sampling's 461.098× P99 comes from completed 31b runs of 3.86–3.88 seconds against an 8.394-ms Full median, not from a timeout.

4.2 Exact-cardinality calibration

Offline, a Datalevin-native factorized evaluator computed exact counts for every connected entity-star subset and every indexed-link input, prior to target-local filtering. Because all 113 JOB entity-star factor graphs are acyclic, it streams star-local query factors and combines their multiplicities by exact tree message passing over shared variables; link inputs combine prefix multiplicities with exact AV-index point counts. This covers the 65,450 logical states and 103,961 distinct link-input count forms without materializing candidate subplans. The checkpointed offline pass completed in under one day.


Exact substitutes these counts without changing enumeration, pruning, operators, or costs. We compare only execution times in 30 paired randomized passes over all 113 queries (3,390 pairs), excluding preparation; all results matched and no run timed out. Both bootstraps resample queries, then seeds: the frozen ablation uses 2,000 replicates, while this later single contrast uses 10,000 to reduce Monte Carlo error.

Full's median paired execution-time ratio to Exact was 1.086× [1.044, 1.160]. Aggregating queries by their median paired ratio

Table 2: Layer-by-layer ablation. C_{run} counts completed trials with $S \geq 10$, C_{TO} counts timeouts, and $C = C_{\text{run}} + C_{\text{TO}}$.

Condition	C_{run}	C_{TO}	C	P95	P99
Full	4	0	4	1.425	2.883
Shrinkage	8	0	8	1.447	2.636
Raw	48	0	48	2.859	12.256
No sampling	176	80	256	93.455	461.098
No counts	127	30	157	6.780	30.059

Note: P95 and P99 are slowdown multiples relative to each condition’s corresponding Full median. No sampling uses 10 randomized schedules per query (1,130 trials); all other conditions use 30 seeds per query (3,390 trials).

**Figure 3: Complementary cumulative distribution function (CCDF) of slowdown relative to the corresponding Full median. Timeouts use the 30-second limit; TO denotes timeout.**

and applying the original JOB study’s buckets, 51.3% of Full queries were within 0.9–1.1 $\times$ of Exact, 20.4% had at least 2 $\times$ slowdown, and none reached 10 $\times$.

For scale, in the same JOB benchmark, PostgreSQL with primary- and foreign-key indexes had 40% of queries at least 2 $\times$ slower than true-cardinality planning; that configuration also disabled non-index nested-loop joins and enabled runtime hash-table resizing [8]. Engines and protocols differ.

Six queries ran over 10% faster under Full, so Exact is not a runtime lower bound: the shared cost model can misrank wall-clock time and favor another plan. Yet with search and costs fixed, the aggregate Full/Exact gap isolates estimation headroom, and the dimensionless weights rank most JOB alternatives usefully once counts are exact.

4.3 Layer-by-layer ablation

Figure 3 shows No sampling with the broadest tail, followed by No counts; Raw separates from Shrinkage and Full around 10 $\times$. Because No sampling has fewer trials, its raw total understates the contrast: 256 of 1,130 trials (22.7%) were catastrophes, versus 157 of 3,390 (4.63%) for No counts. The two failure sets overlap on five queries (8c, 19a, 19c, 25b, and 33b); No sampling nevertheless affected 31 queries versus ten and produced the broader tail.

Full’s four threshold crossings occurred on 23b, 27a, 29b, and 33b. The recorded structural plan recurred in 26 of 30 trials for 23b and all 30 trials for each other query. After excluding the crossing, each same-plan median was within 3.8% of its original Full median; none of the four measurements followed a worker restart. The stable plans make sample-induced plan failure unlikely, but do not identify the physical cause: the harness did not record GC, thermal, SSD, or scheduler counters. We retain all four without attributing them.

Counts prevent severe base-cardinality errors. No counts caused 157 catastrophes across ten queries; query 8c accounted for all 30 timeouts. Of the 127 completed catastrophes, 125 changed operator or traversal order. The 8c timeout concentration is consistent with a misranked starting star, but Table 3 shows that the result does not depend on 8c: dropping it leaves 127 catastrophes across nine queries, including 30 completed catastrophes each on 6b, 17c, and 25b.

Query-local sampling prevents a correlation-and-fanout tail. No sampling produced 256 catastrophes across 31 queries; eight timed out in every schedule (80 timeouts). The timeout tail is conservative because each timeout contributes only the 30-second limit. Excluding those eight all-timeout queries, 176 completed catastrophes remain across 23 queries; dropping any single query leaves at least 246 catastrophes. The fixed plan differed from Full in all ten runs for 99 of 105 queries, indicating systematic rather than timing noise. Counts still provide marginal sizes, so the ablation jointly removes within-star correlation and directional fanout. The durable average can then describe the wrong subpopulation and make expansion appear safe too early.

Regularization suppresses the sampling tail. Raw produced 48 catastrophes, 47 with a plan different from Full. Shrinkage produced eight, five with a different plan, cutting P99 from 12.256 $\times$ to 2.636 $\times$ —a paired reduction of 9.620 [2.693, 11.288]. The catastrophe-rate reduction was 1.180 percentage points, but its query-clustered confidence interval [0.000, 3.187] includes zero. Storage evidence therefore sharply reduces the continuous tail, though the clustered catastrophe-rate interval is inconclusive.

Because Raw and Shrinkage see identical sampled ranks, their difference comes from interpretation rather than observations. Noisy means cross discrete cost boundaries; optimistic fanout errors can move large expansions ahead of filters, whereas moderate overestimates more often merely delay work.

The envelope and shrinkage make different tradeoffs. The envelope produced four observed catastrophes versus eight for Shrinkage, but Full minus Shrinkage in catastrophe rate was -0.118 percentage points [$-0.442, 0.147$]. Shrinkage had lower observed P99, but Full minus Shrinkage was 0.247 [$-0.690, 1.574$]. The envelope guards against estimates that fall below both observation and durable average. Its simpler computation and lower observed catastrophe count motivate the production choice.

In the census, departures from each query-condition modal plan occurred in 15.37% of Raw trials, 11.40% of Shrinkage trials, and 5.79% of Full trials. Thus storage-aware policies reduced seed sensitivity, with the envelope most stable; the metric does not measure agreement with Full.

Table 3: Per-query catastrophe concentration. Q_C : affected queries; $C_{-1} = C - \max_q C_q$: count after dropping the largest contributor. Superscript TO: every trial timed out; zero-count queries omitted.

Condition	Q_C	C	$\max_q C_q$	C_{-1}	Nonzero C_q by query
Full	4	4	1	3	23b, 27a, 29b, 33b (1 each)
Shrinkage	6	8	2	6	22d, 31b (2 each); 1a, 1d, 14b, 21b (1 each)
Raw	8	48	30	18	22d (30); 11d (8); 31b (4); 33c (2); 21a, 24b, 30b, 33a (1 each)
No sampling	31	256	10	246	7b, 8c ^{TO} , 8d ^{TO} , 10c ^{TO} , 12a, 15a, 15c ^{TO} , 15d, 18c, 19c, 22c, 22d, 23a ^{TO} , 23c ^{TO} , 24b, 29a, 29b, 30b ^{TO} , 30c ^{TO} , 31a, 31b, 31c, 33b (10 each); 19a (8); 25c (6); 12c (4); 30a (3); 2b (2); 1c, 6a, 25b (1 each)
No counts	10	157	30	127	6b, 17c, 25b (30 each); 8c ^{TO} (30); 19a (15); 19c (13); 24a (6); 10b, 21c, 33b (1 each)

5 RELATED WORK

Datalevin applies dynamic programming over connected subgraphs after entity-star collapse [11, 13].

Datomic has related Datalog syntax but no cost-based join-order planner: except for a few documented always-beneficial rewrites, it evaluates clauses in source order [5].

A row-store histogram is not the corresponding triple-store baseline. RDF-3X found one-dimensional subject/predicate/object histograms inaccurate on its two more complex datasets. Its effective alternative uses six order-specific histograms with nine join-partner counts per bucket plus frequent-path statistics [12]. Virtuoso likewise reports that generic SQL statistics lose usefulness when RDF occupies one large table and instead samples the live index during optimization [6]. Both systems tie their evidence to triple indexes or graph structure.

Index-based join sampling has been evaluated on JOB [9]; learned estimators target correlated joins [7]. Datalevin instead uses query-time estimates without training or external model state, and our ablations isolate counts, sampling, and regularization.

Robust optimization targets predictable performance under uncertainty [1]. Datalevin derives its uncertainty evidence directly from the indexes used for execution.

6 CONCLUSION

EAV flexibility need not preclude predictable cost-based planning. Index counts prevent base-cardinality errors, query-local samples reveal conditional selectivity and directional fanout, and prior regularization keeps noise from crossing damaging plan boundaries. Hiding counts caused 30 timeouts; disabling sampling caused 80; and shrinkage lowered Raw's P99 from 12.26× to 2.64× on identical ranks. With exact cardinalities, the same optimizer's median paired Full/Exact execution-time ratio was 1.086× [1.044, 1.160], exposing remaining estimator headroom.

The production envelope was simplest and had the lowest observed catastrophe count, though neither its catastrophe-rate nor P99 difference from shrinkage was significant. One model, machine, and analytical workload limit generalization; finite deterministic seeds characterize observed tail risk, not a universal failure probability. Future work should focus sampling near plan-order thresholds and uncertain deeper joins.

REFERENCES

- [1] Brian Babcock and Surajit Chaudhuri. 2005. Towards a Robust Query Optimizer: A Principled and Practical Approach. In *Proceedings of the 2005 ACM SIGMOD International Conference on Management of Data*. 119–130. <https://doi.org/10.1145/1066157.1066172>
- [2] Andreas Brodt, Oliver Schiller, and Bernhard Mitschang. 2011. Efficient Resource Attribute Retrieval in RDF Triple Stores. In *Proceedings of the 20th ACM International Conference on Information and Knowledge Management*. 1445–1454. <https://doi.org/10.1145/2063576.2063785>
- [3] Howard Chu. 2011. MDB: A Memory-Mapped Database and Backend for OpenLDAP. In *Proceedings of the 3rd International Conference on LDAP*. <https://www.openldap.org/pub/hyc/mdb-paper.pdf>
- [4] Datalevin contributors. n.d.. Datalevin: A Simple, Fast, and Versatile Datalog Database. <https://github.com/datalevin/datalevin>. Accessed August 3, 2026.
- [5] Datomic. n.d.. Query Stats. <https://docs.datomic.com/reference/query-stats.html>. Accessed July 31, 2026.
- [6] Orri Erling. n.d.. Implementing a SPARQL-Compliant RDF Triple Store Using a SQL-ORDBMS. OpenLink Virtuoso white paper. Accessed July 31, 2026.
- [7] Andreas Kipf, Thomas Kipf, Bernhard Radke, Viktor Leis, Peter Boncz, and Alfons Kemper. 2019. Learned Cardinalities: Estimating Correlated Joins with Deep Learning. In *9th Biennial Conference on Innovative Data Systems Research*. <https://www.cidrdb.org/cidr2019/papers/p101-kipf-cidr19.pdf>
- [8] Viktor Leis, Andrey Gubichev, Atanas Mirchev, Peter Boncz, Alfons Kemper, and Thomas Neumann. 2015. How Good Are Query Optimizers, Really? *Proceedings of the VLDB Endowment* 9, 3 (2015), 204–215. <https://doi.org/10.14778/2850583.2850594>
- [9] Viktor Leis, Bernhard Radke, Andrey Gubichev, Alfons Kemper, and Thomas Neumann. 2017. Cardinality Estimation Done Right: Index-Based Join Sampling. In *8th Biennial Conference on Innovative Data Systems Research*. <https://www.cidrdb.org/cidr2017/papers/p9-leis-cidr17.pdf>
- [10] David Maier, K. Tuncay Tekle, Michael Kifer, and David S. Warren. 2018. Datalog: Concepts, History, and Outlook. In *Declarative Logic Programming: Theory, Systems, and Applications*. ACM and Morgan & Claypool, 3–100. <https://doi.org/10.1145/3191315.3191317>
- [11] Guido Moerkotte and Thomas Neumann. 2008. Dynamic Programming Strikes Back. In *Proceedings of the 2008 ACM SIGMOD International Conference on Management of Data*. 539–552. <https://doi.org/10.1145/1376616.1376672>
- [12] Thomas Neumann and Gerhard Weikum. 2010. The RDF-3X Engine for Scalable Management of RDF Data. *The VLDB Journal* 19, 1 (2010), 91–113. <https://doi.org/10.1007/s00778-009-0165-y>
- [13] Patricia G. Selinger, Morton M. Astrahan, Donald D. Chamberlin, Raymond A. Lorie, and Thomas G. Price. 1979. Access Path Selection in a Relational Database Management System. In *Proceedings of the 1979 ACM SIGMOD International Conference on Management of Data*. 23–34. <https://doi.org/10.1145/582095.582099>
- [14] Jeffrey Scott Vitter. 1985. Random Sampling with a Reservoir. *ACM Trans. Math. Software* 11, 1 (1985), 37–57. <https://doi.org/10.1145/3147.3165>